

AudioProject

February 19, 2019

0.1 1. Importing needed libraries

```
In [1]: import librosa
import librosa.display

import keras
from keras import regularizers
from keras.models import Sequential, load_model
from keras.layers import Conv2D, MaxPooling2D, Activation, Flatten, Dropout, Dense, GL
from keras.callbacks import ModelCheckpoint, EarlyStopping

import random
import numpy as np
from sklearn.model_selection import train_test_split
import pandas as pd
import matplotlib.pyplot as plt

import fnmatch
import os
import sys

import warnings
warnings.filterwarnings('ignore')

%matplotlib inline
```

Using TensorFlow backend.

```
In [2]: plt.rcParams['figure.figsize'] = [15, 10] # Make output graphs larger in Jupyter Note
```

GPU activation for Nvidia GPUs only

```
In [3]: import tensorflow as tf
sess = tf.Session(config=tf.ConfigProto(log_device_placement=True))
```

```
In [4]: # Check that there is an output
        from keras import backend as k
        k.tensorflow_backend._get_available_gpus()
```

```
Out[4]: []
```

0.2 2. Data Preparation

```
In [4]: data = pd.read_csv('data/metadata.csv')
        data.head(10)
```

```
Out[4]:
```

	slice_file_name	fsID	start	end	salience	fold	classID	\
0	100032-3-0-0.wav	100032	0.000000	0.317551	1	5	3	
1	100263-2-0-117.wav	100263	58.500000	62.500000	1	5	2	
2	100263-2-0-121.wav	100263	60.500000	64.500000	1	5	2	
3	100263-2-0-126.wav	100263	63.000000	67.000000	1	5	2	
4	100263-2-0-137.wav	100263	68.500000	72.500000	1	5	2	
5	100263-2-0-143.wav	100263	71.500000	75.500000	1	5	2	
6	100263-2-0-161.wav	100263	80.500000	84.500000	1	5	2	
7	100263-2-0-3.wav	100263	1.500000	5.500000	1	5	2	
8	100263-2-0-36.wav	100263	18.000000	22.000000	1	5	2	
9	100648-1-0-0.wav	100648	4.823402	5.471927	2	10	1	

	class
0	dog_bark
1	children_playing
2	children_playing
3	children_playing
4	children_playing
5	children_playing
6	children_playing
7	children_playing
8	children_playing
9	car_horn

```
In [5]: print("The original data contains {0} audio files.".format(data.shape[0]))

        # Filter data less than 3 seconds
        dataorg = data[['slice_file_name', 'fold', 'classID', 'class']]
        data3s = data[['slice_file_name', 'fold', 'classID', 'class']][ data['end']-data['start']>3]

        print("The data we will work with contains {0} audio files of 3 seconds length.".format(data3s.shape[0]))

        data3s.head(10)
```

The original data contains 8732 audio files.

The data we will work with contains 7468 audio files of 3 seconds length.

```
Out [5]:
```

	slice_file_name	fold	classID	class
1	100263-2-0-117.wav	5	2	children_playing
2	100263-2-0-121.wav	5	2	children_playing
3	100263-2-0-126.wav	5	2	children_playing
4	100263-2-0-137.wav	5	2	children_playing
5	100263-2-0-143.wav	5	2	children_playing
6	100263-2-0-161.wav	5	2	children_playing
7	100263-2-0-3.wav	5	2	children_playing
8	100263-2-0-36.wav	5	2	children_playing
14	100652-3-0-0.wav	2	3	dog_bark
15	100652-3-0-1.wav	2	3	dog_bark

Different Classes are: 0. air_conditioner

1. car_horn
2. children_playing
3. dog_bark
4. drilling
5. engine_idling
6. gun_shot
7. jackhammer
8. siren
9. street_music

Iterate over all samples in *valid*. For every sample, construct the (128,128) spectrogram

```
In [6]: data3s['path'] = 'fold' + data3s['fold'].astype('str') + '/' + data3s['slice_file_name']
```

```
In [7]: dataset = []
```

```
for row in data3s.itertuples():
    y, sr = librosa.load('data/UrbanSound8K/' + row.path, duration=2.97)
    ps = librosa.feature.melspectrogram(y=y, sr=sr)
    if ps.shape != (128, 128): continue
    dataset.append( (ps, row.classID) )
```

```
In [15]: random.shuffle(dataset)
#X_train, X_test, y_train, y_test = train_test_split(dataset, test_size = .2)
```

```
train = dataset[:7000]
test = dataset[7000:]
```

```
X_train, y_train = zip(*train)
X_test, y_test = zip(*test)
```

```
# Reshape for CNN input
```

```
X_train = np.array([x.reshape( (128, 128, 1) ) for x in X_train])
X_test = np.array([x.reshape( (128, 128, 1) ) for x in X_test])
```

```
# One-Hot encoding for classes
```

```
y_train = np.array(keras.utils.to_categorical(y_train, 10))
y_test = np.array(keras.utils.to_categorical(y_test, 10))
```

```
In [31]: model = keras.models.Sequential()
input_shape=(128, 128, 1)
```

```

# model.add(Conv2D(96, (3, 3), activation='relu', padding = 'same', input_shape=input_shape))
# model.add(Dropout(0.2))
# model.add(Conv2D(96, (3, 3), activation='relu', padding = 'same'))
# model.add(Conv2D(96, (3, 3), activation='relu', padding = 'same', strides = 2))
# model.add(Dropout(0.5))
# model.add(Conv2D(192, (3, 3), activation='relu', padding = 'same'))
# model.add(Conv2D(192, (3, 3), activation='relu', padding = 'same'))
# model.add(Conv2D(256, (3, 3), activation='relu', padding = 'same', strides = 2))
# model.add(Conv2D(512, (3, 3), activation='relu', padding = 'same', strides = 2))
# model.add(Conv2D(512, (3, 3), activation='relu', padding = 'same', strides = 2))
# model.add(Dropout(0.5))
# model.add(Activation('relu'))

```

```

model.add(Conv2D(24, (5, 5), strides=(1, 1), input_shape=input_shape))
model.add(MaxPooling2D((4, 2), strides=(4, 2)))
model.add(Activation('relu'))

```

```

model.add(Conv2D(48, (5, 5), padding="valid"))
model.add(MaxPooling2D((4, 2), strides=(4, 2)))
model.add(Activation('relu'))

```

```

model.add(Conv2D(48, (5, 5), padding="valid"))
model.add(Activation('relu'))

```

```

model.add(Flatten())
model.add(Dropout(rate=0.5))

```

```

model.add(Dense(64))
model.add(Activation('relu'))
model.add(Dropout(rate=0.5))

```

```

model.add(Dense(10))
model.add(Activation('softmax'))

```

```

In [32]: model.compile(
        optimizer="Adam",
        loss="categorical_crossentropy",
        metrics=['accuracy'])

```

```

In [34]: best_model = ModelCheckpoint('best_model.hdf5', save_best_only=True, monitor='val_loss')

```

```

history = model.fit(
    x=X_train,
    y=y_train,

```

```

epochs=80,
batch_size=500,
validation_data= (X_test, y_test),
callbacks=[best_model] )

plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

# Plot training & validation loss values
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Train on 7000 samples, validate on 467 samples

```

Epoch 1/80
7000/7000 [=====] - 64s 9ms/step - loss: 3.1472 - acc: 0.1744 - val_loss: 3.1472
Epoch 2/80
7000/7000 [=====] - 62s 9ms/step - loss: 2.2972 - acc: 0.2214 - val_loss: 2.2972
Epoch 3/80
7000/7000 [=====] - 63s 9ms/step - loss: 2.1687 - acc: 0.2643 - val_loss: 2.1687
Epoch 4/80
7000/7000 [=====] - 65s 9ms/step - loss: 2.0756 - acc: 0.2973 - val_loss: 2.0756
Epoch 5/80
7000/7000 [=====] - 62s 9ms/step - loss: 1.9899 - acc: 0.3291 - val_loss: 1.9899
Epoch 6/80
7000/7000 [=====] - 63s 9ms/step - loss: 1.8871 - acc: 0.3426 - val_loss: 1.8871
Epoch 7/80
7000/7000 [=====] - 64s 9ms/step - loss: 1.8010 - acc: 0.3650 - val_loss: 1.8010
Epoch 8/80
7000/7000 [=====] - 65s 9ms/step - loss: 1.7353 - acc: 0.3871 - val_loss: 1.7353
Epoch 9/80
7000/7000 [=====] - 63s 9ms/step - loss: 1.6994 - acc: 0.4079 - val_loss: 1.6994
Epoch 10/80
7000/7000 [=====] - 63s 9ms/step - loss: 1.6176 - acc: 0.4234 - val_loss: 1.6176
Epoch 11/80
7000/7000 [=====] - 61s 9ms/step - loss: 1.5471 - acc: 0.4491 - val_loss: 1.5471

```



Epoch 12/80
7000/7000 [=====] - 61s 9ms/step - loss: 1.4719 - acc: 0.4697 - val_loss: 1.5000
Epoch 13/80
7000/7000 [=====] - 63s 9ms/step - loss: 1.4027 - acc: 0.4996 - val_loss: 1.4500
Epoch 14/80
7000/7000 [=====] - 63s 9ms/step - loss: 1.4062 - acc: 0.5049 - val_loss: 1.4500
Epoch 15/80
7000/7000 [=====] - 61s 9ms/step - loss: 1.3408 - acc: 0.5203 - val_loss: 1.4500
Epoch 16/80
7000/7000 [=====] - 62s 9ms/step - loss: 1.2855 - acc: 0.5466 - val_loss: 1.4500
Epoch 17/80
7000/7000 [=====] - 63s 9ms/step - loss: 1.2133 - acc: 0.5711 - val_loss: 1.4500
Epoch 18/80
7000/7000 [=====] - 61s 9ms/step - loss: 1.1491 - acc: 0.5923 - val_loss: 1.4500
Epoch 19/80
7000/7000 [=====] - 62s 9ms/step - loss: 1.1265 - acc: 0.6120 - val_loss: 1.4500
Epoch 20/80
7000/7000 [=====] - 63s 9ms/step - loss: 1.1108 - acc: 0.6147 - val_loss: 1.4500
Epoch 21/80
7000/7000 [=====] - 63s 9ms/step - loss: 1.0957 - acc: 0.6230 - val_loss: 1.4500
Epoch 22/80
7000/7000 [=====] - 62s 9ms/step - loss: 1.1544 - acc: 0.5963 - val_loss: 1.4500
Epoch 23/80
7000/7000 [=====] - 64s 9ms/step - loss: 1.0688 - acc: 0.6279 - val_loss: 1.4500
Epoch 24/80
7000/7000 [=====] - 63s 9ms/step - loss: 1.0161 - acc: 0.6506 - val_loss: 1.4500
Epoch 25/80
7000/7000 [=====] - 62s 9ms/step - loss: 0.9760 - acc: 0.6657 - val_loss: 1.4500
Epoch 26/80
7000/7000 [=====] - 63s 9ms/step - loss: 0.9507 - acc: 0.6859 - val_loss: 1.4500
Epoch 27/80
7000/7000 [=====] - 64s 9ms/step - loss: 0.9194 - acc: 0.6881 - val_loss: 1.4500
Epoch 28/80
7000/7000 [=====] - 61s 9ms/step - loss: 0.8916 - acc: 0.7041 - val_loss: 1.4500
Epoch 29/80
7000/7000 [=====] - 61s 9ms/step - loss: 0.8624 - acc: 0.7150 - val_loss: 1.4500
Epoch 30/80
7000/7000 [=====] - 63s 9ms/step - loss: 0.8203 - acc: 0.7254 - val_loss: 1.4500
Epoch 31/80
7000/7000 [=====] - 61s 9ms/step - loss: 0.8061 - acc: 0.7359 - val_loss: 1.4500
Epoch 32/80
7000/7000 [=====] - 63s 9ms/step - loss: 0.7826 - acc: 0.7377 - val_loss: 1.4500
Epoch 33/80
7000/7000 [=====] - 63s 9ms/step - loss: 0.7930 - acc: 0.7394 - val_loss: 1.4500
Epoch 34/80
7000/7000 [=====] - 62s 9ms/step - loss: 0.7724 - acc: 0.7456 - val_loss: 1.4500
Epoch 35/80
7000/7000 [=====] - 62s 9ms/step - loss: 0.7616 - acc: 0.7463 - val_loss: 1.4500

Epoch 36/80
7000/7000 [=====] - 62s 9ms/step - loss: 0.7378 - acc: 0.7550 - val_loss: 0.7550
Epoch 37/80
7000/7000 [=====] - 62s 9ms/step - loss: 0.7452 - acc: 0.7550 - val_loss: 0.7550
Epoch 38/80
7000/7000 [=====] - 63s 9ms/step - loss: 0.7373 - acc: 0.7596 - val_loss: 0.7596
Epoch 39/80
7000/7000 [=====] - 62s 9ms/step - loss: 0.7463 - acc: 0.7560 - val_loss: 0.7560
Epoch 40/80
7000/7000 [=====] - 61s 9ms/step - loss: 0.6755 - acc: 0.7781 - val_loss: 0.7781
Epoch 41/80
7000/7000 [=====] - 63s 9ms/step - loss: 0.6580 - acc: 0.7859 - val_loss: 0.7859
Epoch 42/80
7000/7000 [=====] - 62s 9ms/step - loss: 0.6318 - acc: 0.7871 - val_loss: 0.7871
Epoch 43/80
7000/7000 [=====] - 65s 9ms/step - loss: 0.6105 - acc: 0.7999 - val_loss: 0.7999
Epoch 44/80
7000/7000 [=====] - 65s 9ms/step - loss: 0.6134 - acc: 0.7964 - val_loss: 0.7964
Epoch 45/80
7000/7000 [=====] - 63s 9ms/step - loss: 0.5847 - acc: 0.8056 - val_loss: 0.8056
Epoch 46/80
7000/7000 [=====] - 62s 9ms/step - loss: 0.6409 - acc: 0.7901 - val_loss: 0.7901
Epoch 47/80
7000/7000 [=====] - 64s 9ms/step - loss: 0.6897 - acc: 0.7791 - val_loss: 0.7791
Epoch 48/80
7000/7000 [=====] - 62s 9ms/step - loss: 0.6602 - acc: 0.7873 - val_loss: 0.7873
Epoch 49/80
7000/7000 [=====] - 60s 9ms/step - loss: 0.6149 - acc: 0.8020 - val_loss: 0.8020
Epoch 50/80
7000/7000 [=====] - 61s 9ms/step - loss: 0.5794 - acc: 0.8124 - val_loss: 0.8124
Epoch 51/80
7000/7000 [=====] - 60s 9ms/step - loss: 0.5647 - acc: 0.8207 - val_loss: 0.8207
Epoch 52/80
7000/7000 [=====] - 60s 9ms/step - loss: 0.5271 - acc: 0.8241 - val_loss: 0.8241
Epoch 53/80
7000/7000 [=====] - 60s 9ms/step - loss: 0.5163 - acc: 0.8287 - val_loss: 0.8287
Epoch 54/80
7000/7000 [=====] - 61s 9ms/step - loss: 0.5293 - acc: 0.8219 - val_loss: 0.8219
Epoch 55/80
7000/7000 [=====] - 61s 9ms/step - loss: 0.5279 - acc: 0.8260 - val_loss: 0.8260
Epoch 56/80
7000/7000 [=====] - 63s 9ms/step - loss: 0.5070 - acc: 0.8346 - val_loss: 0.8346
Epoch 57/80
7000/7000 [=====] - 62s 9ms/step - loss: 0.5107 - acc: 0.8387 - val_loss: 0.8387
Epoch 58/80
7000/7000 [=====] - 67s 10ms/step - loss: 0.4924 - acc: 0.8421 - val_loss: 0.8421
Epoch 59/80
7000/7000 [=====] - 62s 9ms/step - loss: 0.4655 - acc: 0.8489 - val_loss: 0.8489

```

Epoch 60/80
7000/7000 [=====] - 62s 9ms/step - loss: 0.4708 - acc: 0.8420 - val_loss: 0.4708
Epoch 61/80
7000/7000 [=====] - 63s 9ms/step - loss: 0.4517 - acc: 0.8514 - val_loss: 0.4517
Epoch 62/80
7000/7000 [=====] - 63s 9ms/step - loss: 0.4530 - acc: 0.8486 - val_loss: 0.4530
Epoch 63/80
7000/7000 [=====] - 63s 9ms/step - loss: 0.4285 - acc: 0.8631 - val_loss: 0.4285
Epoch 64/80
7000/7000 [=====] - 61s 9ms/step - loss: 0.4259 - acc: 0.8614 - val_loss: 0.4259
Epoch 65/80
7000/7000 [=====] - 63s 9ms/step - loss: 0.4447 - acc: 0.8583 - val_loss: 0.4447
Epoch 66/80
7000/7000 [=====] - 62s 9ms/step - loss: 0.4347 - acc: 0.8541 - val_loss: 0.4347
Epoch 67/80
7000/7000 [=====] - 62s 9ms/step - loss: 0.4380 - acc: 0.8597 - val_loss: 0.4380
Epoch 68/80
7000/7000 [=====] - 63s 9ms/step - loss: 0.4893 - acc: 0.8496 - val_loss: 0.4893
Epoch 69/80
7000/7000 [=====] - 62s 9ms/step - loss: 0.5084 - acc: 0.8354 - val_loss: 0.5084
Epoch 70/80
7000/7000 [=====] - 62s 9ms/step - loss: 0.4822 - acc: 0.8431 - val_loss: 0.4822
Epoch 71/80
7000/7000 [=====] - 64s 9ms/step - loss: 0.4742 - acc: 0.8450 - val_loss: 0.4742
Epoch 72/80
7000/7000 [=====] - 70s 10ms/step - loss: 0.4376 - acc: 0.8577 - val_loss: 0.4376
Epoch 73/80
7000/7000 [=====] - 72s 10ms/step - loss: 0.4144 - acc: 0.8689 - val_loss: 0.4144
Epoch 74/80
7000/7000 [=====] - 73s 10ms/step - loss: 0.5492 - acc: 0.8267 - val_loss: 0.5492
Epoch 75/80
7000/7000 [=====] - 70s 10ms/step - loss: 0.5343 - acc: 0.8346 - val_loss: 0.5343
Epoch 76/80
7000/7000 [=====] - 65s 9ms/step - loss: 0.4617 - acc: 0.8466 - val_loss: 0.4617
Epoch 77/80
7000/7000 [=====] - 62s 9ms/step - loss: 0.4304 - acc: 0.8623 - val_loss: 0.4304
Epoch 78/80
7000/7000 [=====] - 63s 9ms/step - loss: 0.4022 - acc: 0.8686 - val_loss: 0.4022
Epoch 79/80
7000/7000 [=====] - 64s 9ms/step - loss: 0.4233 - acc: 0.8604 - val_loss: 0.4233
Epoch 80/80
7000/7000 [=====] - 65s 9ms/step - loss: 0.3835 - acc: 0.8736 - val_loss: 0.3835

```

Test loss: 1.2244046683995575
Test accuracy: 0.5952890787185897

```
In [5]: best_model = load_model('best_model.hdf5')
```

```
In [26]: !rm data/myAudioFiles/.DS_Store
```

rm: data/myAudioFiles/.DS_Store: No such file or directory

```
In [25]: path = 'data/myAudioFiles/'
```

```
for audio_file in sorted (os.listdir(path)):
    y, sr = librosa.load(path + audio_file, duration=2.97)
    ps = librosa.feature.melspectrogram(y=y, sr=sr)

    k = np.array(ps) #seven
    y = k.reshape(1, 128, 128, 1)
    # predict
    print('File \t{0} \tpredicted as class {1} \t{2}'.format(audio_file, best_model.predict_classes(y)[0] == int(audio_file[0]))) # predict class

# plt.rcParams['figure.figsize'] = [3, 2]
# plt.plot(np.arange(10), best_model.predict(y).ravel()) # prediction likelihood
# plt.show()
```

File	0.air_condition.m4a	predicted as class [0]	True
File	1.car_horn_short.m4a	predicted as class [1]	True
File	2.children_playing.m4a	predicted as class [2]	True
File	4.drilling_machine.m4a	predicted as class [4]	True
File	6.gunshot_firing.m4a	predicted as class [8]	False
File	7.jackhammer_tool.m4a	predicted as class [7]	True
File	8.siren_alarm.m4a	predicted as class [8]	True
File	9.street_music.m4a	predicted as class [9]	True

```
In [21]: plt.rcParams['figure.figsize'] = [15, 10] # Make output graphs larger in Jupyter Notebook
```

```
In [22]: !rm data/gun_firing/.DS_Store
```

rm: data/gun_firing/.DS_Store: No such file or directory

```
In [24]: path = 'data/gun_firing/'
```

```
for audio_file in sorted (os.listdir(path)):
    print(path + audio_file)
```

```

y, sr = librosa.load(path + audio_file, duration=2.97)
ps = librosa.feature.melspectrogram(y=y, sr=sr)
k = np.array(ps)

plt.figure(audio_file)
plt.title(audio_file)
plt.imshow(k)
plt.plot()

librosa.display.specshow(ps, y_axis='mel', x_axis='time')

data/gun_firing/6.gunshot_firing.m4a
data/gun_firing/train_sample.m4a

```

